



REGRESSION ANALYSIS (ASSESSMENT - 1)

CSE4020(MACHINE LEARNING)LAB:L49-L50



**JANUARY 22, 2022
ANIRUDH VADERA
20BCE2940**

QUESTION:

1. Linear Regression:

Read the students' height and weight from the csv file. Understand the relationship between the features using Linear Regression. Predict the value of weight for the given height of a student. Find the performance of the model using appropriate metrics. Visualize the result in scatter plot. Divide the dataset into training and testing. Predict the value for the test set.

Expected Output

1. Initial scatter plot
2. Scatter plot along with regression line
3. List the x value and the predicted y value
4. Display the error

2. Multi Linear Regression

Multi Linear Regression Predict the value of house price given the features about the house. Use Multi Linear Regression to predict the value. Divide the data set into train and test with 65% and 35% respectively. Visualize the result in scatter plot and display the error


Expected Output

- 1) Scatter plot with regression line
- 2) List the x value and the predicted y value
- 3) Display the error

→ Linear Regression

Description:

Linear regression is an approach to model the relationship between a single **dependent variable** (target variable) and one (simple regression) or more (multiple regression) **independent variables**. The linear regression model assumes a linear relationship between the input and output variables. If this relationship is present, we can estimate the coefficients required by the model to make predictions on new data.



Formula Used:

Hypothesis function for Linear Regression:

$$Y(\text{Pred}) = b_0 + b_1 \cdot X$$

b₀ = Intercept

b₁ = Coefficient

$$\text{Error} = \sum_{i=1}^n (\text{actual_output} - \text{predicted_output}) ** 2$$

Figure 2: Error Calculation

$$b_0 = \bar{y} - b_1 \bar{x}$$

Figure 3: Intercept Calculation

$$b_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

Figure 4: Co-efficient Formula

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
from sklearn.linear_model import LinearRegression  
  
from sklearn.metrics import mean_absolute_error  
  
from sklearn.metrics import mean_squared_error  
  
from sklearn import metrics
```

Importing the dataset

```
df=pd.read_csv("C:/Users/Anirudh/OneDrive/Desktop/weight-height.csv")  
  
print("The dataset is as following : [10000 rows x 3 columns]")  
  
print(df)  
  
print("\n")
```

Check for missing values

```
print("Checking for missing values :")  
  
print(df.isnull().sum())  
  
print("\n")
```

Printing the header of the dataset

```
print("Dataset Header : ")  
  
print(df.head())  
  
print("\n")
```

Information regarding the columns

```
print("Information regarding the columns : ")  
  
print("Height : inches")  
  
print("Weight : pounds")  
  
print(df.info())  
  
print("\n")
```

Information related to the dataset

```
print("Dataset Details : ")  
  
print(df.describe())
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
print("\n")
```

```
# Plot Height vs Weight.....
```

```
# Blue points - Male
```

```
# Magenta points - Female
```

```
ax1 = df[df['Gender'] == 'Male'].plot(kind='scatter', x='Height', y='Weight', color='blue', alpha=0.5,  
figsize=(10, 7))
```

```
df[df['Gender'] == 'Female'].plot(kind='scatter', x='Height', y='Weight', color='magenta', alpha=0.5,  
figsize=(10, 7), ax=ax1)
```

```
plt.legend(labels=['Males', 'Females'])
```

```
plt.title('Relationship between Height and Weight', size=24)
```

```
plt.xlabel('Height (inches)', size=18)
```

```
plt.ylabel('Weight (pounds)', size=18)
```

```
# Fitting a Regression Model
```

```
# 20% of dataset - Testing
```

```
# 80% of dataset - Training
```

```
# Random state = 123
```

```
X = df[["Height"]]
```

```
Y = df[["Weight"]]
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=123)
```

```
lin_reg = LinearRegression()
```

```
lin_reg.fit(X_train, Y_train)
```

```
# Predicting the test set values
```

```
lin_pred = lin_reg.predict(X_test)
```

```
print("The X Values are : ")
```

```
print(X_test)
```

```
print("\n")
```

```
print("The predicted Y_Values are : ")
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
print(lin_pred)
print("\n")
```

Visualizing the result in form of a scatterplot (Training Set)

```
plt.scatter(X_train, Y_train, color = 'Magenta')
plt.plot(X_train, lin_reg.predict(X_train), color = 'blue')
plt.legend(labels=['Height vs Weight plot(Training Set)', 'Regression Line'])
plt.title('Relationship between Height and Weight using Regression Line(Training Set)', size=24)
plt.xlabel('Height (inches)', size=18)
plt.ylabel('Weight (pounds)', size=18)
```

Visualizing the result in form of a scatterplot (Testing Set)

```
plt.scatter(X_test, Y_test, color = 'Magenta')
plt.plot(X_train, lin_reg.predict(X_train), color = 'blue')
plt.legend(labels=['Height vs Weight plot(Test Set)', 'Regression Line'])
plt.title('Relationship between Height and Weight using Regression Line(Test Set)', size=24)
plt.xlabel('Height (inches)', size=18)
plt.ylabel('Weight (pounds)', size=18)
```

Checking the accuracy of our model

```
print("Accuracy of the model : ")
print("\n")
print('Mean Absolute Error:', mean_absolute_error(Y_test, lin_pred))
print('Mean Squared Error:', mean_squared_error(Y_test, lin_pred))
print('Mean Root Squared Error:', np.sqrt(mean_squared_error(Y_test, lin_pred)))
print("\n")
print('Variance score: %.2f' % lin_reg.score(X_test, Y_test))
print('Coefficients: ', lin_reg.coef_)
print('R square = ', metrics.r2_score(Y_test, lin_pred)*100)
```

Making new predictions based on given height

ANIRUDH VADERA REGRESSION ANALYSIS (ASSESSMENT - 1)

```
given_height = int(input("Enter the height of the student whose weight you want to predict (in inches) : "));
```

```
print("\n")
```

```
predict_weight = lin_reg.predict([[given_height]])
```

```
print("The predicted weight is (in pounds) : ", predict_weight)
```

The Linear Regression Equation

```
print("The Coefficient is : ", lin_reg.coef_)
```

```
print("The Intercept is : ", lin_reg.intercept_)
```

```
print("The Linear regression line is : Y = ", lin_reg.intercept_[0], "+", lin_reg.coef_[0][0], "*X")
```

Residual Analysis

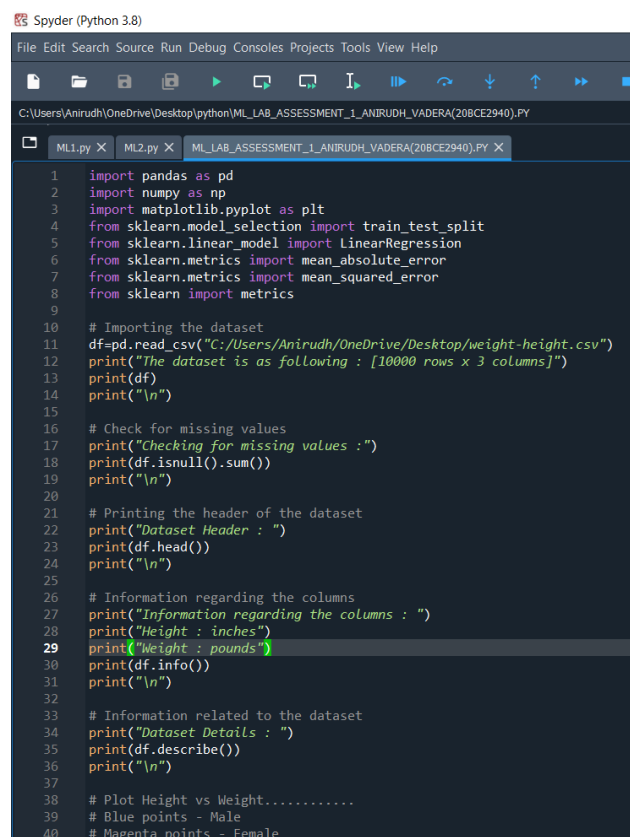
```
plt.scatter(Y_test, lin_pred, color="Magenta")
```

```
plt.title('Checking difference between the predicted Y_Values and the original Y_Values', size=24)
```

```
plt.xlabel('Y_test set', size=18)
```

```
plt.ylabel('Y Predicted', size=18)
```

Code Snippets:



```
Spyder (Python 3.8)
File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Users\Anirudh\OneDrive\Desktop\python\ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(208CE2940).PY

ML1.py X ML2.py X ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(208CE2940).PY X

1  import pandas as pd
2  import numpy as np
3  import matplotlib.pyplot as plt
4  from sklearn.model_selection import train_test_split
5  from sklearn.linear_model import LinearRegression
6  from sklearn.metrics import mean_absolute_error
7  from sklearn.metrics import mean_squared_error
8  from sklearn import metrics
9
10 # Importing the dataset
11 df=pd.read_csv("C:/Users/Anirudh/OneDrive/Desktop/weight-height.csv")
12 print("The dataset is as following : [10000 rows x 3 columns]")
13 print(df)
14 print("\n")
15
16 # Check for missing values
17 print("Checking for missing values :")
18 print(df.isnull().sum())
19 print("\n")
20
21 # Printing the header of the dataset
22 print("Dataset Header : ")
23 print(df.head())
24 print("\n")
25
26 # Information regarding the columns
27 print("Information regarding the columns : ")
28 print("Height : inches")
29 print("Weight : pounds")
30 print(df.info())
31 print("\n")
32
33 # Information related to the dataset
34 print("Dataset Details : ")
35 print(df.describe())
36 print("\n")
37
38 # Plot Height vs Weight.....
39 # Blue points - Male
40 # Magenta points - Female
```

Output and Results:

Dataset Details:

Dataset:

```
Python 3.8.12 (default, Oct 12 2021, 03:01:40) [MSC v.1916 64 bit (AMD64)]
Type "copyright", "credits" or "license" for more information.

IPython 7.29.0 -- An enhanced Interactive Python.

In [1]: runfile('C:/Users/Anirudh/OneDrive/Desktop/python/ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940).PY',
OneDrive/Desktop/python')
The dataset is as following : [10000 rows x 3 columns]
  Gender    Height    Weight
0     Male  73.847017  241.893563
1     Male  68.781904  162.310473
2     Male  74.110105  212.740856
3     Male  71.730978  220.042470
4     Male  69.881796  206.349801
...
9995  Female  66.172652  136.777454
9996  Female  67.067155  170.867906
9997  Female  63.867992  128.475319
9998  Female  69.034243  163.852461
9999  Female  61.944246  113.649103

[10000 rows x 3 columns]

Checking for missing values :
Gender    0
Height    0
Weight    0
dtype: int64
```

As the missing values is none we can proceed further:

Dataset Details:

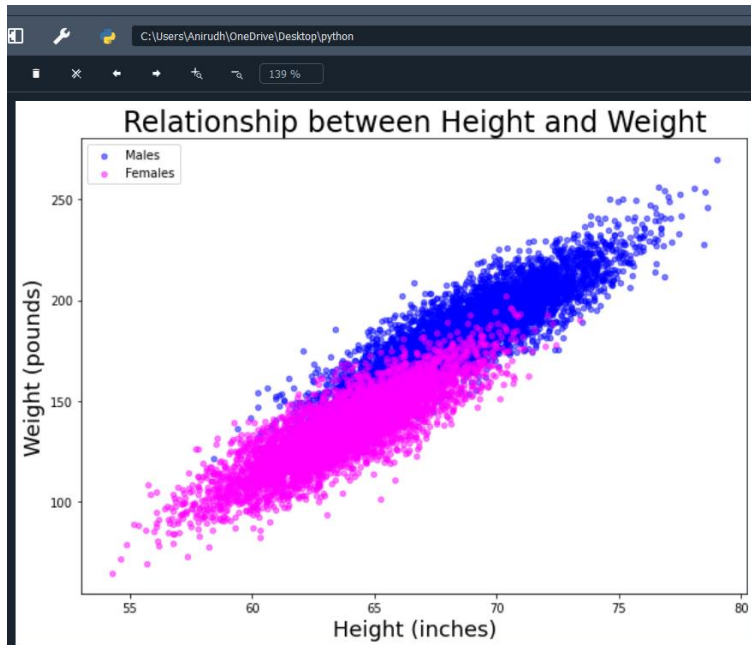
```
Dataset Header :
  Gender    Height    Weight
0     Male  73.847017  241.893563
1     Male  68.781904  162.310473
2     Male  74.110105  212.740856
3     Male  71.730978  220.042470
4     Male  69.881796  206.349801

Information regarding the columns :
Height : inches
Weight : pounds
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 3 columns):
#   Column    Non-Null Count  Dtype
---  ---
0   Gender    10000 non-null  object
1   Height    10000 non-null  float64
2   Weight    10000 non-null  float64
dtypes: float64(2), object(1)
memory usage: 234.5+ KB
None
```

```
Dataset Details :
              Height              Weight
count  10000.000000  10000.000000
mean      66.367560    161.440357
std       3.847528     32.108439
min       54.263133     64.700127
25%       63.505620    135.818051
50%       66.318070    161.212928
75%       69.174262    187.169525
max       78.998742    269.989699
```


Initial Scatterplot:

ScatterPlot between weight and height to understand the relationship between the two:



The X_Values and Predicted Y_Values:

Fitting the Regression Model:

80% data for training and 20% for testing:

The X-Values for which Y-Values to be predicted are:

```
The X Values are :
      Height
2656  73.181767
445   71.433376
9505  60.026750
332   66.556587
4168  71.035509
...
8018  63.992565
6463  64.388396
2883  65.625006
7895  61.530110
620   65.892753

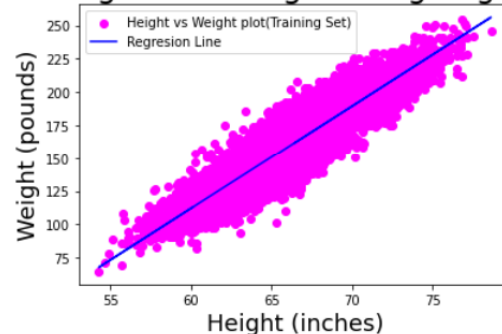
[2000 rows x 1 columns]

The predicted Y_Values are :
[[214.02231091]
 [200.50905719]
 [112.34767127]
 ...
 [155.61638129]
 [123.96708917]
 [157.68578703]]
```

Regression Results:

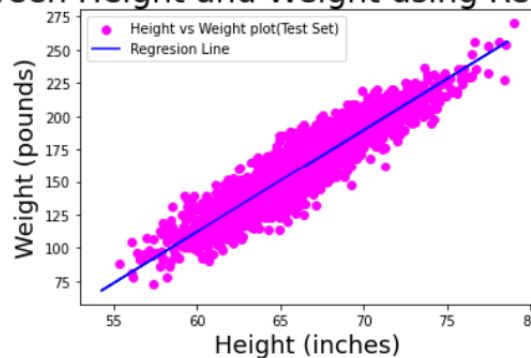
ScatterPlot for training set:

Relationship between Height and Weight using Regression Line(Training Set)



ScatterPlot for testing set:

Relationship between Height and Weight using Regression Line(Test Set)



Accuracy Analysis(Errors):

Checking the Accuracy of the model:

Accuracy of the model :

Mean Absolute Error: 9.688834054963015
Mean Squared Error: 143.2255601011165
Mean Root Squared Error: 11.9676881686112

Variance score: 0.86
Coefficients: [[7.72896259]]
R square = 86.49031737206691

Predicting the weight based on given height:

Enter the height of the student whose weight you want to predict (in inches) : 170

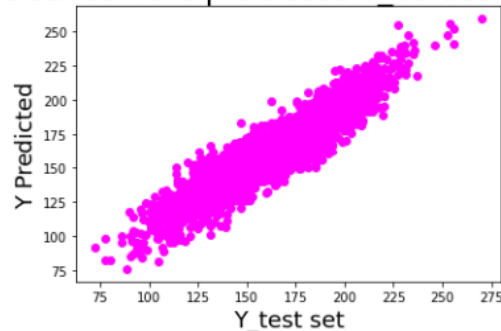
The predicted weight is (in pounds) : [[962.32680767]]

The Regression Equation:

```
The Coefficient is : [[7.72896259]]  
The Intercept is : [-351.59683192]  
The Linear regression line is :  $Y = -351.59683191643967 + 7.728962585782046 * X$ 
```

The Difference between the predicted and original Y-values:

Checking difference between the predicted Y Values and the original Y_Values



We get a straight line with positive slope

Inference:

1. Positive Linear Relationship

The initial ScatterPlot shows that on increasing the height the weight also increases

2. We can see from the output below that R square is 0.864 that is 86.4% which is good and mean squared error is a bit high at 143 but manageable.

3. The variance in predicted y_values and the actual y_values is not so much which can be seen through the value of mean absolute error which is 9.688

4. The Regression Equation is:

$$Y = -351.59683191643967 + 7.728962585782046 * X$$


Y = weight, X = height

5. As we can see in the last scatterplot (the difference between predicted and actual y_values) we get almost a slanting straight line in positive direction (and positive slope) which shows values were well predicted.

→ Multi Linear Regression

Description:

Linear regression is an approach to model the relationship between a single **dependent variable** (target variable) and one (simple regression) or more (multiple regression) **independent variables**. The linear regression model assumes a linear relationship between the input and output variables. If this relationship is present, we can estimate the coefficients required by the model to make predictions on new data.



Formula Used:

Hypothesis function for Linear Regression:

$$Y(\text{Pred}) = b_0 + b_1 \cdot X_1 + b_2 \cdot X_2 \dots b_n \cdot X_n$$

b0 = Intercept

b1 = Coefficient of X1

b2 = Coefficient of X2

bn = Coefficient of Xn

$$b_0 = \bar{y} - b_1 \bar{x}$$

Figure 3: Intercept Calculation

$$\text{Error} = \sum_{i=1}^n (\text{actual_output} - \text{predicted_output}) ** 2$$

Figure 2: Error Calculation

$$b_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

Figure 4: Co-efficient Formula

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
from sklearn.metrics import mean_absolute_error  
from sklearn.metrics import mean_squared_error  
from sklearn import metrics  
import seaborn as sns
```

Importing the dataset

```
df=pd.read_csv("C:/Users/Anirudh/OneDrive/Desktop/Housing.csv")  
print("The dataset is as following : [545 rows x 13 columns]")  
print(df)  
print("\n")
```

Check for missing values

```
print("Checking for missing values :")  
print(df.isnull().sum())  
print("\n")
```

Printing the header of the dataset

```
print("Dataset Header : ")  
print(df.head())  
print("\n")
```

Information regarding the columns

```
print("Information regarding the columns : ")  
print("Height : inches")  
print("Weight : pounds")  
print(df.info())  
print("\n")
```

Information related to the dataset

```
print("Dataset Details : ")  
print(df.describe())
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
print("\n")
```

Data Preparation

You can see that your dataset has many columns with values as 'Yes' or 'No'.

But in order to fit a regression line, we would need numerical values and not string.

Hence, we need to convert them to 1s and 0s, where 1 is a 'Yes' and 0 is a 'No'.

List of variables to map

```
varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea']
```

Defining the map function

```
def binary_map(x):  
    return x.map({'yes': 1, "no": 0})
```

Applying the function to the housing list

```
df[varlist] = df[varlist].apply(binary_map)
```

Dummy Variables

The variable furnishingstatus has three levels. We need to convert these levels into integer as well.

For this, we will use something called dummy variables.

Get the dummy variables for the feature 'furnishingstatus' and store it in a new variable - 'status'

```
status = pd.get_dummies(df['furnishingstatus'])
```

Now, you don't need three columns.

You can drop the furnished column, as the type of furnishing can be identified with just the last two columns where —

00 will correspond to furnished

01 will correspond to unfurnished

10 will correspond to semi-furnished

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

Let's drop the first column from status df using 'drop_first = True'

```
status = pd.get_dummies(df['furnishingstatus'], drop_first = True)
```

Add the results to the original housing dataframe

```
df = pd.concat([df, status], axis = 1)
```

Drop 'furnishingstatus' as we have created the dummies for it

```
df.drop(['furnishingstatus'], axis = 1, inplace = True)
```

Now let's see the head of our dataframe.

```
print("After Trimming and correcting the dataset looks like follows : ")
```

```
print(df.head())
```

Fitting a Regression Model

35% of dataset - Testing

65% of dataset - Training

Random state = 10

```
df_train, df_test = train_test_split(df, test_size=0.35, train_size=0.65, random_state=10)
```

Let's check the correlation coefficients to see which variables are highly correlated

```
plt.figure(figsize = (16, 10))
```

```
sns.heatmap(df_train.corr(), annot = True, cmap="YlGnBu")
```

```
plt.show()
```

We can see there are some factors that affects the price of a house to great extent

These columns are :

Area

Bathrooms

Stories

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

Air conditioning

This means changing value of even one of these factors will greatly affect the pricing of the house

Showing the relationship between the price and the most dominant factor that is area

```
plt.scatter(df["area"], df["price"], color = 'Magenta')  
plt.legend(labels=['Area vs Height plot'])  
plt.title('Relationship between Price and area of the houses', size=24)  
plt.xlabel('Area', size=18)  
plt.ylabel('Price', size=18)
```

Dividing into X and Y sets for the model building

```
Y_train = df_train.pop('price')  
X_train = df_train  
# Using the Linear Regression Model  
lin_reg = LinearRegression()  
lin_reg.fit(X_train, Y_train)
```

```
Y_test = df_test.pop("price")  
X_test = df_test
```

Predicting the test set values

```
lin_pred = lin_reg.predict(X_test)
```

```
print("The X Values are : ")
```

```
print(X_test)
```

```
print("\n")
```

```
print("The predicted Y_Values are : ")
```


ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
print(lin_pred)

print("\n")

print("The values look something like this : ")

X_test["Price"] = lin_pred

print(X_test)

X_test.drop(["Price"],axis=1,inplace=True)

new_lin_reg = LinearRegression()

new_lin_reg.fit(X_train["area"].values.reshape(-1,1), Y_train)
```

Visualizing the result in form of a scatterplot (Training Set)

We are only visualizing for the area factor as it is the most dominant

```
plt.scatter(X_train["area"], Y_train, color = 'Magenta')

plt.plot(X_train["area"], new_lin_reg.predict(X_train["area"].values.reshape(-1,1)), color = 'blue')

plt.legend(labels=['Area vs Price plot(Training Set)', 'Regresion Line'])

plt.title('Relationship between Area and Price of Houses using Regression Line(Training Set)',
size=24)

plt.xlabel('Area' , size=18)

plt.ylabel('Price', size=18)
```

Visualizing the result in form of a scatterplot (Testing Set)

We are only visualizing for the area factor as it is the most dominant

```
plt.scatter(X_test["area"], Y_test, color = 'Magenta')

plt.plot(X_train["area"], new_lin_reg.predict(X_train["area"].values.reshape(-1,1)), color = 'blue')

plt.legend(labels=['Area vs Price plot(Training Set)', 'Regresion Line'])

plt.title('Relationship between Area and Price of Houses using Regression Line(Test Set)', size=24)

plt.xlabel('Area' , size=18)

plt.ylabel('Price', size=18)
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

Checking the accuracy of our model

```
print("Accuracy of the model : ")
print("\n")
print('Mean Absolute Error:', mean_absolute_error(Y_test, lin_pred))
print('Mean Squared Error:', mean_squared_error(Y_test, lin_pred))
print('Mean Root Squared Error:', np.sqrt(mean_squared_error(Y_test, lin_pred)))
print("\n")
print('Variance score: %.2f' % lin_reg.score(X_test, Y_test))
print('R square = ', metrics.r2_score(Y_test, lin_pred)*100)
```

Plot the histogram of the error terms

```
fig = plt.figure()
sns.distplot((Y_test - lin_pred), bins = 20)
fig.suptitle('Error Terms', fontsize = 20)          # Plot heading
plt.xlabel('Errors', fontsize = 18)                 # X-label
```

Making new predictions based on given features of the house

```
print("Enter the features of the house : ")
area = int(input("Enter the area of the house : "))
bedrooms = int(input("Enter the no of bedrooms in the house : "))
bathrooms = int(input("Enter the no of bathrooms in the house : "))
stories = int(input("Enter the stories in the house : "))
mainroad = int(input("Is there a attached mainroad 'Yes':1 'No':0 "))
guestroom = int(input("Is there a attached guestroom 'Yes':1 'No':0 "))
basement = int(input("Is there a basement 'Yes':1 'No':0 "))
hotwaterheating = int(input("Is there hot water heating 'Yes':1 'No':0 "))
airconditioning = int(input("Is there air conditioning 'Yes':1 'No':0 "))
parking = int(input("Is there parking nearby 'Yes':1 'No':0 "))
prefarea = int(input("Is there prefarea nearby 'Yes':1 'No':0 "))
furnished = int(input("Is the house semi-furnished:10 furnished:00 or unfurnished:01 :: "))
if (furnished==0):
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
semi_furnished = 0
unfurnished = 0
elif(furnished==1):
    semi_furnished = 0
    unfurnished = 1
else:
    semi_furnished = 1
    unfurnished = 0

predict_price =
lin_reg.predict([[area,bedrooms,bathrooms,stories,mainroad,guestroom,basement,hotwaterheating
,airconditioning,parking,prefarea,semi_furnished,unfurnished]])

print("The predicted price of the house is : ", predict_price)

# The Linear Regression Equation

# Checking the coefficients

print("The Coefficient are : ")

coeff_df = pd.DataFrame(lin_reg.coef_,X_train.columns,columns=["Coefficient"])

print(coeff_df)

print("The Intercept is : ", lin_reg.intercept_)

print("The Linear regression line is : Price = ", lin_reg.intercept_ , "+" , lin_reg.coef_[0] , "*area" , "+"
, lin_reg.coef_[1] , "*bedrooms" , "+" , lin_reg.coef_[2] , "*bathrooms" , "+" , lin_reg.coef_[3] ,
"*stories" , "+" , lin_reg.coef_[4] , "*mainroad" , "+" , lin_reg.coef_[5] , "*guestroom" , "+" ,
lin_reg.coef_[6] , "*basement" , "+" , lin_reg.coef_[7] , "*hotwaterheating" , "+" , lin_reg.coef_[8] ,
"*airconditioning" , "+" , lin_reg.coef_[9] , "*parking" , "+" , lin_reg.coef_[10] , "*prefarea")

# Residual Ananlysis

plt.scatter(Y_test,lin_pred,color="Magenta")

plt.title('Checking difference between the predicted Y_Values and the original Y_Values', size=24)

plt.xlabel('Y_test set', size=18)

plt.ylabel('Y Predicted', size=18)

# Plot the histogram of the error terms

fig = plt.figure()

sns.distplot((Y_test - lin_pred), bins = 20)

fig.suptitle('Error Terms', fontsize = 20)          # Plot heading

plt.xlabel('Errors', fontsize = 18)
```

ANIRUDH VADERA REGRESSION ANALYSIS (ASSESSMENT - 1)

Code Snippets:

```
Spyder (Python 3.8)
File Edit Search Source Run Debug Consoles Projects Tools View Help

...\\Anirudh\\OneDrive\\Desktop\\python\\ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940)\\MULTIPLE_REGRESSION.PY

Client.py X Server_.py X ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940).PY X ML_LAB_ASSESSMENT_1_A

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_absolute_error
7 from sklearn.metrics import mean_squared_error
8 from sklearn import metrics
9 import seaborn as sns
10
11 # Importing the dataset
12 df=pd.read_csv("C:/Users/Anirudh/OneDrive/Desktop/Housing.csv")
13 print("The dataset is as following : [545 rows x 13 columns]")
14 print(df)
15 print("\n")
16
17 # Check for missing values
18 print("Checking for missing values :")
19 print(df.isnull().sum())
20 print("\n")
21
22 # Printing the header of the dataset
23 print("Dataset Header : ")
24 print(df.head())
25 print("\n")
26
27 # Information regarding the columns
28 print("Information regarding the columns : ")
29 print("Height : inches")
30 print("Weight : pounds")
31 print(df.info())
32 print("\n")
33
34 # Information related to the dataset
35 print("Dataset Details : ")
36 print(df.describe())
37 print("\n")
38
```

```
Spyder (Python 3.8)
File Edit Search Source Run Debug Consoles Projects Tools View Help

C:\Users\Anirudh\OneDrive\Desktop\py

...\\Anirudh\\OneDrive\\Desktop\\python\\ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940)\\MULTIPLE_REGRESSION.PY

Client.py X Server_.py X ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940).PY X ML_LAB_ASSESSMENT_1_ANIRUDH_VADERA(20BCE2940)MULTIPLE_REGRESSION.PY X

39 # Data Preparation
40 # You can see that your dataset has many columns with values as 'Yes' or 'No'.
41 # But in order to fit a regression line, we would need numerical values and not string.
42 # Hence, we need to convert them to 1s and 0s, where 1 is a 'Yes' and 0 is a 'No'.
43
44 # List of variables to map
45 varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea']
46
47 # Defining the map function
48 def binary_map(x):
49     return x.map({'yes': 1, "no": 0})
50
51 # Applying the function to the housing list
52 df[varlist] = df[varlist].apply(binary_map)
53
54 # Dummy Variables
55 # The variable furnishingstatus has three levels. We need to convert these levels into integer as well.
56 # For this, we will use something called dummy variables.
57 # Get the dummy variables for the feature 'furnishingstatus' and store it in a new variable - 'status'
58 status = pd.get_dummies(df['furnishingstatus'])
59
60 # Now, you don't need three columns.
61 # You can drop the furnished column, as the type of furnishing can be identified with just the last two columns where
62
63 # 00 will correspond to furnished
64 # 01 will correspond to unfurnished
65 # 10 will correspond to semi-furnished
66
67 # Let's drop the first column from status df using 'drop_first = True'
68
69 status = pd.get_dummies(df['furnishingstatus'], drop_first = True)
70
71 # Add the results to the original housing dataframe
72 df = pd.concat([df, status], axis = 1)
73
74 # Drop 'furnishingstatus' as we have created the dummies for it
75 df.drop(['furnishingstatus'], axis = 1, inplace = True)
76 # Now let's see the head of our dataframe.
77 print(df.head())
```

Output and Results:

Dataset Details:

Dataset:

```
In [1]: runfile('C:/Users/Anirudh/OneDrive/Desktop/python/ML_LAB_ASSESSMENT_1_
wdir='C:/Users/Anirudh/OneDrive/Desktop/python')
The dataset is as following : [545 rows x 13 columns]
   price  area  bedrooms  ...  parking  prefarea  furnishingstatus
0  13300000  7420         4  ...      2         yes      furnished
1  12250000  8960         4  ...      3         no      furnished
2  12250000  9960         3  ...      2         yes  semi-furnished
3  12215000  7500         4  ...      3         yes      furnished
4  11410000  7420         4  ...      2         no      furnished
..  ...  ...  ...  ...  ...  ...  ...
540  1820000  3000         2  ...      2         no      unfurnished
541  1767150  2400         3  ...      0         no  semi-furnished
542  1750000  3620         2  ...      0         no      unfurnished
543  1750000  2910         3  ...      0         no      furnished
544  1750000  3850         3  ...      0         no      unfurnished

[545 rows x 13 columns]

Checking for missing values :
price          0
area           0
bedrooms       0
bathrooms      0
stories        0
mainroad       0
guestroom      0
basement       0
hotwaterheating 0
airconditioning 0
parking        0
prefarea       0
furnishingstatus 0
dtype: int64
```

As the missing values is none we can proceed further:

Dataset Details:

```
Dataset Header :
   price  area  bedrooms  ...  parking  prefarea  furnishingstatus
0  13300000  7420         4  ...      2         yes      furnished
1  12250000  8960         4  ...      3         no      furnished
2  12250000  9960         3  ...      2         yes  semi-furnished
3  12215000  7500         4  ...      3         yes      furnished
4  11410000  7420         4  ...      2         no      furnished

[5 rows x 13 columns]

Information regarding the columns :
Height : inches
Weight : pounds
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   price               545 non-null   int64
1   area                545 non-null   int64
2   bedrooms            545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
None
```

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

```
Dataset Details :
      price      area  ...  stories  parking
count  5.450000e+02  545.000000  ...  545.000000  545.000000
mean    4.766729e+06  5150.541284  ...    1.805505    0.693578
std     1.870440e+06  2170.141023  ...    0.867492    0.861586
min     1.750000e+06  1650.000000  ...    1.000000    0.000000
25%     3.430000e+06  3600.000000  ...    1.000000    0.000000
50%     4.340000e+06  4600.000000  ...    2.000000    0.000000
75%     5.740000e+06  6360.000000  ...    2.000000    1.000000
max     1.330000e+07  16200.000000  ...    4.000000    3.000000
```

Data Preparation

You can see that your dataset has many columns with values as 'Yes' or 'No'.

But in order to fit a regression line, we would need numerical values and not string.

Hence, we need to convert them to 1s and 0s, where 1 is a 'Yes' and 0 is a 'No'.

Dummy Variables

The variable furnishingstatus has three levels. We need to convert these levels into integer as well.

For this, we will use something called dummy variables.

Get the dummy variables for the feature 'furnishingstatus' and store it in a new variable - 'status'

Now, you don't need three columns.

You can drop the furnished column, as the type of furnishing can be identified with just the last two columns where —

00 will correspond to furnished

01 will correspond to unfurnished

10 will correspond to semi-furnished

```
After Trimming and correcting the dataset looks like follows :
      price  area  bedrooms  ...  prefarea  semi-furnished  unfurnished
0  13300000  7420         4  ...         1             0             0
1  12250000  8960         4  ...         0             0             0
2  12250000  9960         3  ...         1             1             0
3  12215000  7500         4  ...         1             0             0
4  11410000  7420         4  ...         0             0             0

[5 rows x 14 columns]
```

Initial Scatterplot:

We now draw a heatmap of covariance of each column with one another.

We now get some following observations:

We can see there are some factors that affects the price of a house to great extent

These columns are:


Area

Bathrooms

Stories

Air conditioning


This means changing value of even one of these factors will greatly affect the pricing of the house



As the area is the most dominant factor in affecting the value of price of the house we will draw scatterplots between area and price of the house

ScatterPlot between Price and Area to understand the relationship between the two:

Relationship between Price and area of the houses



The X_Values and Predicted Y_Values:

Fitting the Regression Model:

65% data for training and 35% for testing:

The X-Values for which Y-Values to be predicted are:

The X Values are :								
	area	bedrooms	bathrooms	...	prefarea	semi-furnished	unfurnished	
482	3150	3	1	...	0	0	1	
314	4040	2	1	...	0	1	0	
383	4500	4	2	...	0	0	0	
487	5400	4	1	...	0	0	1	
43	6000	4	2	...	0	1	0	
..	
12	6550	4	2	...	1	1	0	
97	6400	3	1	...	1	1	0	
130	4800	3	1	...	0	0	1	
33	5960	3	3	...	0	0	1	
465	3800	2	1	...	0	0	1	

The predicted Y_Values are :				
[2660628.13495901	3835819.77139557	5820822.60902346	3767930.49734531
	6670701.48650124	6724771.63134602	3545985.06635784	3100031.51306977
	5554503.39233212	3199668.22391934	5047667.16279301	2161341.91419495
	4304172.15652621	4746707.19522931	6730414.60542925	2648241.04780559
	7299096.84535021	6511349.52910361	5182838.03057654	4448705.93984871
	2538223.51269119	1981826.64972116	5923383.65631762	5656007.7591878
	6650998.59277708	6227896.37806669	3819226.88533281	3661500.68911227
	3744646.53737908	3673436.01446116	4320861.53054848	4784148.06135873
	5690613.88966773	5421403.67760399	7422575.60614913	4978800.24562416
	6589798.86011599	2335117.37297983	5139405.48180192	5655270.1751458
	4955655.47660856	5352075.8841257	2853042.3237546	6371473.85858323
	5591239.29939988	7888140.31937221	8179835.08408416	6308371.48751598
	3531006.48625688	5022139.32687507	3422394.49865358	6006108.73703623
	4148960.8385987	4792267.10805539	2815543.64539553	5906169.94684547
	5251100.43535874	3722859.07495565	3843125.0718775	4494856.65418746
	4793195.96418023	4963726.14788397	5303068.43449917	3644597.89959067
	4317494.74840909	6753219.65260247	3047506.81059606	2504775.87533761
	2426800.24508677	7226894.00319624	4146264.70977009	6562335.34613059
	6063259.30206841	3619479.58429875	2780810.44339752	3520494.46129019

ANIRUDH VADERA
REGRESSION ANALYSIS (ASSESSMENT - 1)

The values look something like this :


	area	bedrooms	bathrooms	...	semi-furnished	unfurnished	Price
482	3150	3	1	...	0	1	2.660628e+06
314	4040	2	1	...	1	0	3.835820e+06
383	4500	4	2	...	0	0	5.820823e+06
487	5400	4	1	...	0	1	3.767930e+06
43	6000	4	2	...	1	0	6.670701e+06
..
12	6550	4	2	...	1	0	7.295587e+06
97	6400	3	1	...	1	0	6.450643e+06
130	4800	3	1	...	0	1	3.594404e+06
33	5960	3	3	...	0	1	6.222597e+06
465	3800	2	1	...	0	1	2.580317e+06

[191 rows x 14 columns]

Regression Results:


ScatterPlot for training set:

Relationship between Area and Price of Houses using Regression Line(Training Set)



ScatterPlot for testing set:

Relationship between Area and Price of Houses using Regression Line(Test Set)



Accuracy Analysis(Errors):

Checking the Accuracy of the model:

Accuracy of the model :

Mean Absolute Error: 724444.1496414025
Mean Squared Error: 816243286778.123
Mean Root Squared Error: 903461.834710312

Variance score: 0.74
R square = 73.5848179015316

Predicting the price of the house based on given features:

```
Enter the features of the house :  
  
Enter the area of the house : 3000  
  
Enter the no of bedrooms in the house : 2  
  
Enter the no of bathrooms in the house : 3  
  
Enter the stories in the house : 2  
  
Is there a attached mainroad 'Yes':1 'No':0 1  
  
Is there a attached guestroom 'Yes':1 'No':0 1  
  
Is there a basement 'Yes':1 'No':0 1  
  
Is there hot water heating 'Yes':1 'No':0 1  
  
Is there air conditioning 'Yes':1 'No':0 1  
  
Is there parking nearby 'Yes':1 'No':0 1  
  
Is there prefarea nearby 'Yes':1 'No':0 1  
  
Is the house semi-furnished:10 furnished:00 or unfurnihsed:01 :: 00
```


```
The predicted price of the house is : [8252024.39997437]
```

The Regression Equation:

```
The Coefficient are :  
Coefficient  
area          2.476073e+02  
bedrooms      1.720664e+05  
bathrooms     7.828399e+05  
stories       4.473094e+05  
mainroad      3.781196e+05  
guestroom     1.627630e+05  
basement      4.316503e+05  
hotwaterheating 8.178225e+05  
airconditioning 1.017280e+06  
parking       3.279801e+05  
prefarea      6.940205e+05  
semi-furnished 1.549957e+04  
unfurnished   -4.052882e+05  
The Intercept is : 92295.7479408225  
The Linear regression line is : Price = 92295.7479408225 + 247.60726134315706 *area + 172066.38956123567 *bedrooms +  
782839.8873357589 *bathrooms + 447309.35295425705 *stories + 378119.63480296044 *mainroad + 162762.95492393035 *guestroom +  
431650.2903978061 *basement + 817822.5179267152 *hotwaterheating + 1017279.6925102217 *airconditioning + 327980.1105063436  
*parking + 694020.5198978384 *prefarea
```


The Difference between the predicted and original Y-values:

Checking difference between the predicted Y_Values and the original Y_Values



We get a straight line with positive slope

Error Terms:



We notice that errors are in acceptable range

Inference:

1. Positive Linear Relationship

The initial ScatterPlot shows that on increasing the area and the other constraints the weight also increases

2. We can see from the output below that R square is 0.735 that is ~74% which is good.

3. The variance score is 0.74 which means the predicted values is much closer to the actual values

4. The Regression Equation is:

Price = 92295.7479408225 + 247.60726134315706 *area + 172066.38956123567 *bedrooms + 782839.8873357589 *bathrooms + 447309.35295425705 *stories + 378119.63480296044 *mainroad + 162762.95492393035 *guestroom + 431650.2903978061 *basement + 817822.5179267152 *hotwaterheating + 1017279.6925102217 *airconditioning + 327980.1105063436 *parking + 694020.5198978384 *prefarea

5. As we can see in the last scatterplot (the difference between predicted and actual y_values) we get almost a slanting straight line in positive direction (and positive slope) which shows values were well predicted.

6. The error barplot shows that the error is in acceptable range

7. The area is the most dominant factor in affecting the value of price of the house